

# **Introduction**

- 1. Purpose of Database Systems**
- 2. View of Data**
- 3. Data Models**
- 4. Data Definition Language**
- 5. Data Manipulation Language**
- 6. Transaction Management**
- 7. Storage Management**
- 8. Database Administrator**
- 9. Database Users**
- 10. Overall System Structure**

## **Other database models**

A variety of other database models have been or are still used today.

### **Inverted file model**

A database built with the inverted file structure is designed to facilitate fast full text searches. In this model, data content is indexed as a series of keys in a lookup table, with the values pointing to the location of the associated files. This structure can provide nearly instantaneous reporting in big data and analytics, for instance.

This model has been used by the ADABAS database management system of Software AG since 1970, and it is still supported today.

### **Flat model**

The flat model is the earliest, simplest data model. It simply lists all the data in a single table, consisting of columns and rows. In order to access or manipulate the data, the computer has to read the entire flat file into memory, which makes this model inefficient for all but the smallest data sets.

### **Multidimensional model**

This is a variation of the relational model designed to facilitate improved analytical processing. While the relational model is optimized for online transaction processing (OLTP), this model is designed for online analytical processing (OLAP).

Each cell in a dimensional database contains data about the dimensions tracked by the database. Visually, it's like a collection of cubes, rather than two-dimensional tables.

## **Semistructured model**

In this model, the structural data usually contained in the database schema is embedded with the data itself. Here the distinction between data and schema is vague at best. This model is useful for describing systems, such as certain Web-based data sources, which we treat as databases but cannot constrain with a schema. It's also useful for describing interactions between databases that don't adhere to the same schema.

## **Context model**

This model can incorporate elements from other database models as needed. It cobbles together elements from object-oriented, semistructured, and network models.

## **Associative model**

This model divides all the data points based on whether they describe an entity or an association. In this model, an entity is anything that exists independently, whereas an association is something that only exists in relation to something else.

The associative model structures the data into two sets:

A set of items, each with a unique identifier, a name, and a type

A set of links, each with a unique identifier and the unique identifiers of a source, verb, and target. The stored fact has to do with the source, and each of the three identifiers may refer either to a link or an item.

Other, less common database models include:

Semantic model, which includes information about how the stored data relates to the real world

XML database, which allows data to be specified and even stored in XML format

Named graph

Triplestore

### **NoSQL database models**

In addition to the object database model, other non-SQL models have emerged in contrast to the relational model:


The **graph database model**, which is even more flexible than a network model, allowing any node to connect with any other.

The **multivalue model**, which breaks from the relational model by allowing attributes to contain a list of data rather than a single data point.


The **document model**, which is designed for storing and managing documents or semi-structured data, rather than atomic data.

### **Databases on the Web**

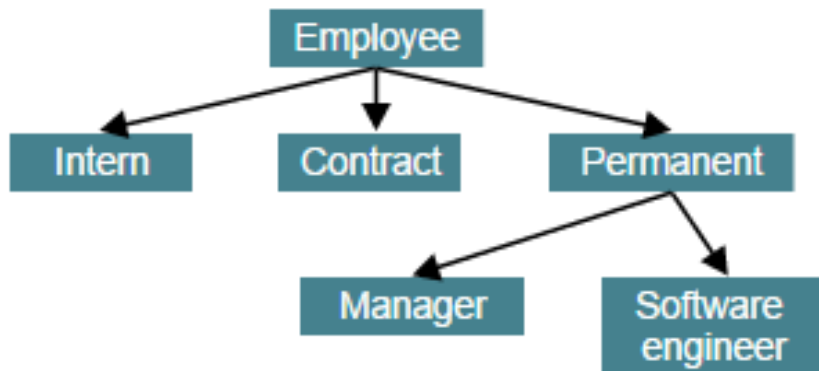
# DBMS Models



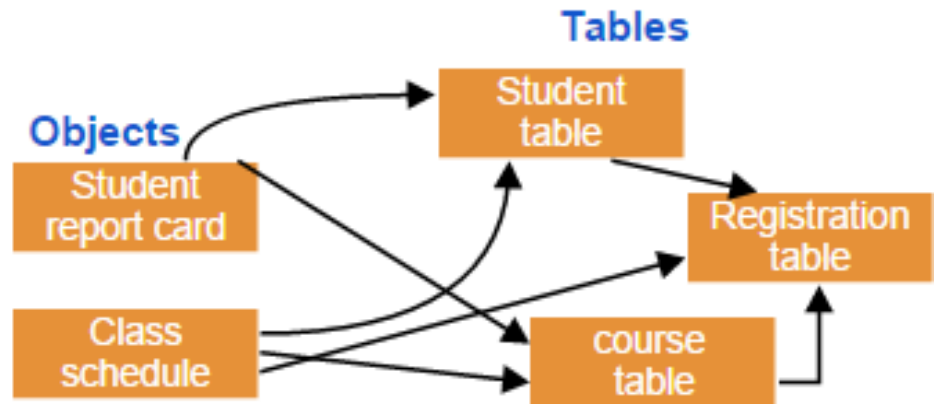
Network Database Model



Relational Database Model



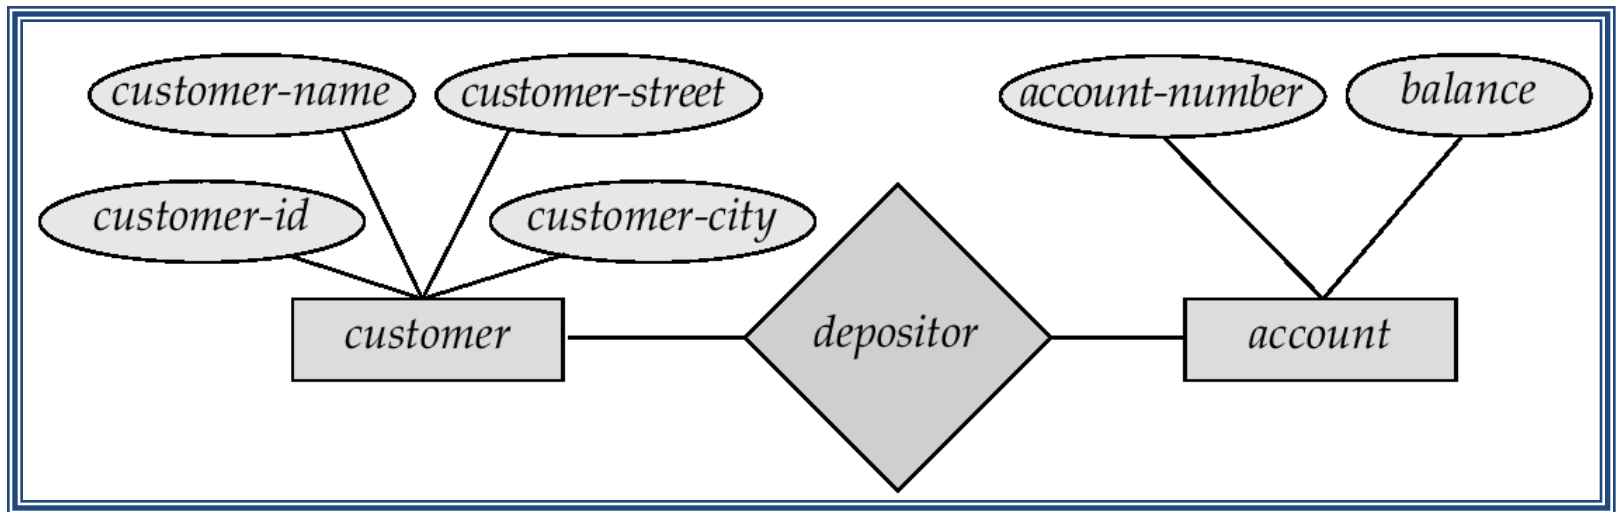
Hierarchical Database Model



Object-oriented Database Model

# Entity-Relationship Model

Example of schema in the entity-relationship model

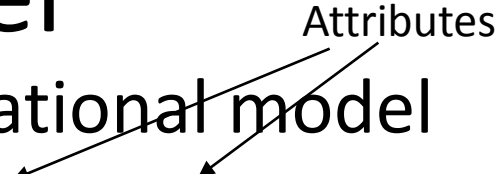


# Entity Relationship Model (Cont.)

- E-R model of real world
  - Entities (objects)
    - E.g. customers, accounts, bank branch
  - Relationships between entities
    - E.g. Account A-101 is held by customer Johnson
    - Relationship set *depositor* associates customers with accounts
- Widely used for database design
  - Database design in E-R model usually converted to design in the relational model (coming up next) which is used for storage and processing

# Relational Model

- Example of tabular data in the relational model



<i>Customer-id</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>	<i>account-number</i>
192-83-7465	Johnson	Alma	Palo Alto	A-101
019-28-3746	Smith	North	Rye	A-215
192-83-7465	Johnson	Alma	Palo Alto	A-201
321-12-3123	Jones	Main	Harrison	A-217
019-28-3746	Smith	North	Rye	A-201



# A Sample Relational Database

<i>customer-id</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
192-83-7465	Johnson	12 Alma St.	Palo Alto
019-28-3746	Smith	4 North St.	Rye
677-89-9011	Hayes	3 Main St.	Harrison
182-73-6091	Turner	123 Putnam Ave.	Stamford
321-12-3123	Jones	100 Main St.	Harrison
336-66-9999	Lindsay	175 Park Ave.	Pittsfield
019-28-3746	Smith	72 North St.	Rye

(a) The *customer* table

<i>account-number</i>	<i>balance</i>
A-101	500
A-215	700
A-102	400
A-305	350
A-201	900
A-217	750
A-222	700

(b) The *account* table

<i>customer-id</i>	<i>account-number</i>
192-83-7465	A-101
192-83-7465	A-201
019-28-3746	A-215
677-89-9011	A-102
182-73-6091	A-305
321-12-3123	A-217
336-66-9999	A-222
019-28-3746	A-201

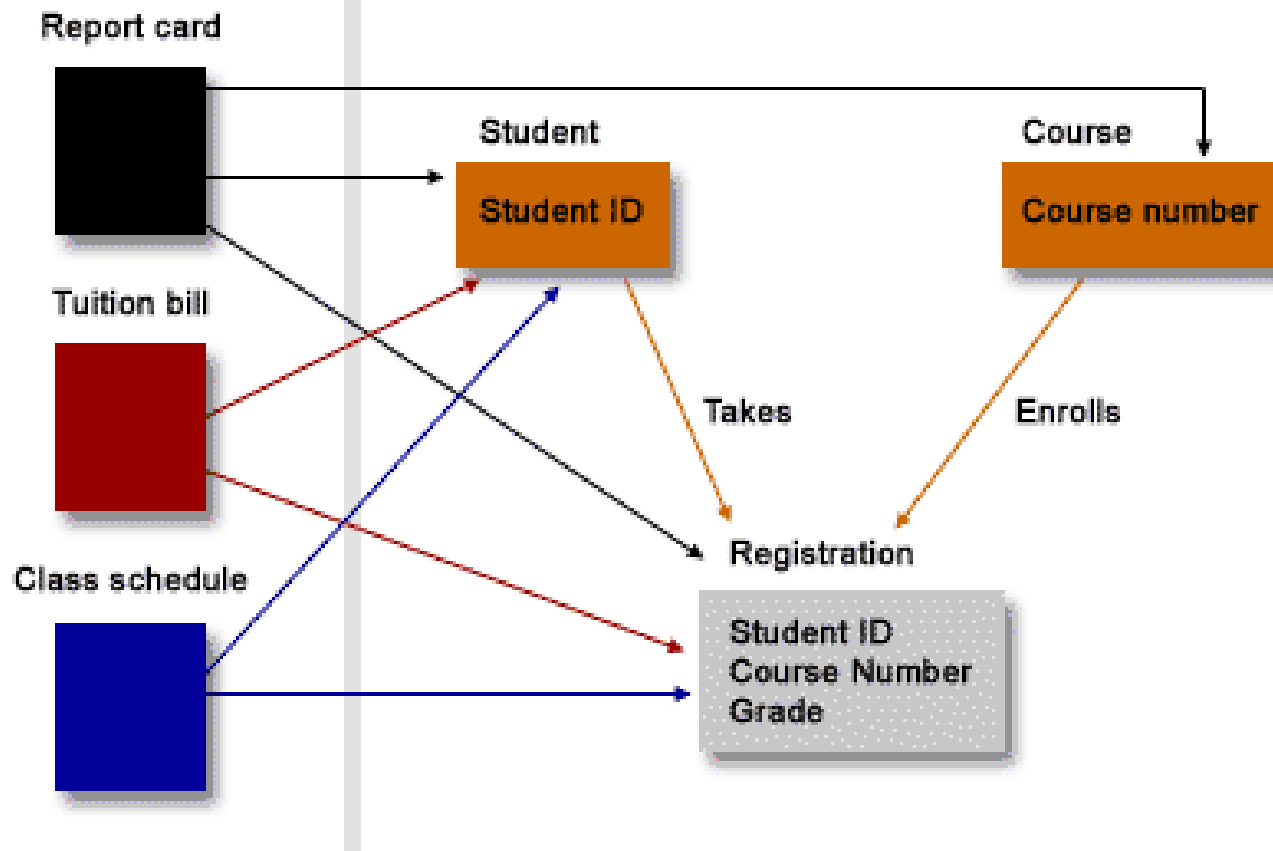
(c) The *depositor* table

# Data Definition Language (DDL)

- Specification notation for defining the database schema
  - E.g.  
**create table** *account* (  
                  *account-number*   **char**(10),  
                  *balance*           **integer**)
- DDL compiler generates a set of tables stored in a *data dictionary*
- Data dictionary contains metadata (i.e., data about data)
  - database schema
  - Data *storage and definition* language
    - language in which the storage structure and access methods used by the database system are specified
    - Usually an extension of the data definition language

## Objects

## Database entries



# Data Manipulation Language (DML)

- Language for accessing and manipulating the data organized by the appropriate data model
  - DML also known as query language
- Two classes of languages
  - Procedural – user specifies what data is required and how to get those data
  - Nonprocedural – user specifies what data is required without specifying how to get those data
- SQL is the most widely used query language

## **DML**

DML is abbreviation of **Data Manipulation Language**. It is used to retrieve, store, modify, delete, insert and update data in database.

Examples: SELECT, UPDATE, INSERT statements

## **DDL**

DDL is abbreviation of **Data Definition Language**. It is used to create and modify the structure of database objects in database.

Examples: CREATE, ALTER, DROP statements

## **DCL**

DCL is abbreviation of **Data Control Language**. It is used to create roles, permissions, and referential integrity as well it is used to control access to database by securing it.

Examples: GRANT, REVOKE statements

## **TCL**

TCL is abbreviation of **Transactional Control Language**. It is used to manage different transactions occurring within a database.

Examples: COMMIT, ROLLBACK statements

A **data manipulation language (DML)** is a family of syntax elements similar to a computer programming language used for selecting, inserting, deleting and updating data in a database. Performing read-only queries of data is sometimes also considered a component of DML.

A data manipulation language (DML) is a family of computer languages including commands permitting users to manipulate data in a database. This manipulation involves inserting data into database tables, retrieving existing data, deleting data from existing tables and modifying existing data. DML is mostly incorporated in SQL databases.

**Data Definition Language (DDL)** is a standard for commands that **define** the different structures in a database. DDL statements create, modify, and remove database objects such as tables, indexes, and users. Common DDL statements are CREATE, ALTER, and DROP.

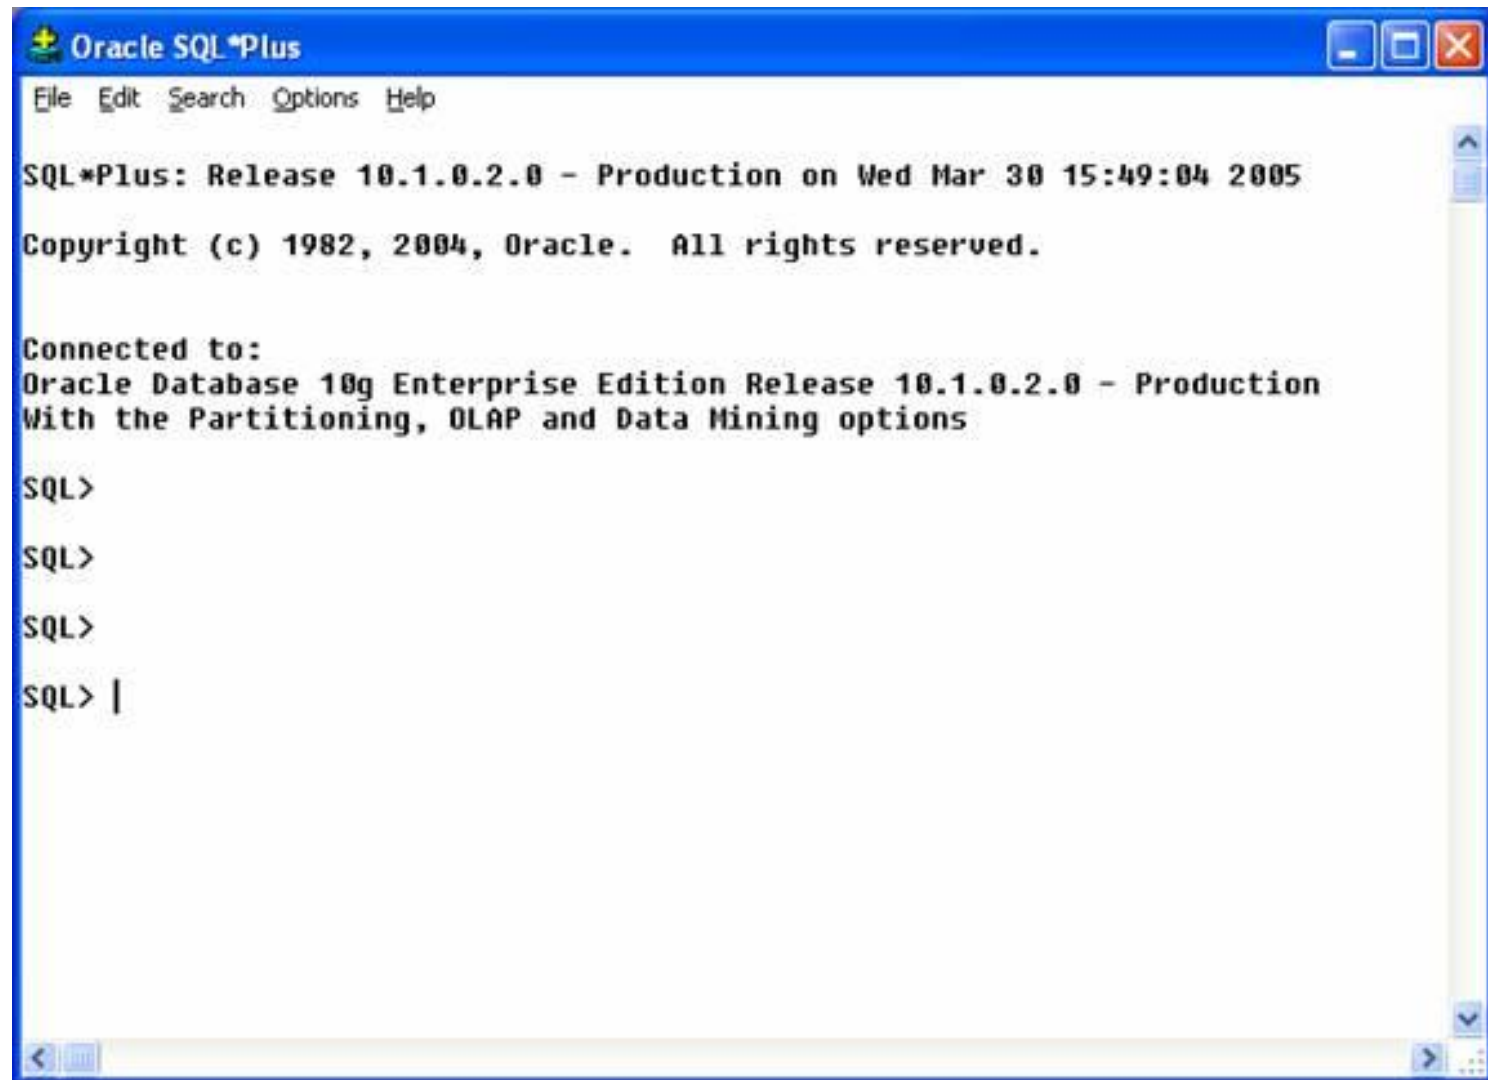
A **data control language (DCL)** is a syntax similar to a computer programming language used to **control** access to **data** stored in a database (Authorization). In particular, it is a component of Structured Query **Language (SQL)**. Examples of DCL commands include: GRANT to allow specified users to perform specified tasks.

# SQL

- SQL: widely used non-procedural language
  - E.g. find the name of the customer with customer-id 192-83-7465

```
select  customer.customer-name
from    customer
where   customer.customer-id = '192-83-7465'
```
  - E.g. find the balances of all accounts held by the customer with customer-id 192-83-7465

```
select  account.balance
from    depositor, account
where   depositor.customer-id = '192-83-7465' and
          depositor.account-number = account.account-
          number
```
- Application programs generally access databases through one of
  - Language extensions to allow embedded SQL
  - Application program interface (e.g. ODBC/JDBC) which allow SQL queries to be sent to a database



The image shows a screenshot of the Oracle SQL\*Plus application window. The title bar at the top reads "Oracle SQL\*Plus" and includes standard window control buttons (minimize, maximize, close). Below the title bar is a menu bar with the following options: File, Edit, Search, Options, and Help. The main text area displays the following information:

```
SQL*Plus: Release 10.1.0.2.0 - Production on Wed Mar 30 15:49:04 2005  
Copyright (c) 1982, 2004, Oracle. All rights reserved.  
  
Connected to:  
Oracle Database 10g Enterprise Edition Release 10.1.0.2.0 - Production  
With the Partitioning, OLAP and Data Mining options  
  
SQL>  
SQL>  
SQL>  
SQL> |
```

The window also features a vertical scrollbar on the right side and a horizontal scrollbar at the bottom.



ssh Packt - ssh oracle@192.168.2.10

SQL> desc EXT\_EMPLOYEES

Name	Null?	Type
EMPLOYEE_ID		NUMBER(6)
FIRST_NAME		VARCHAR2(20)
LAST_NAME		VARCHAR2(25)
EMAIL		VARCHAR2(25)
PHONE_NUMBER		VARCHAR2(20)
HIRE_DATE		DATE
JOB_ID		VARCHAR2(10)
SALARY		NUMBER(8,2)
COMMISSION_PCT		NUMBER(2,2)
MANAGER_ID		NUMBER(6)
DEPARTMENT_ID		NUMBER(4)

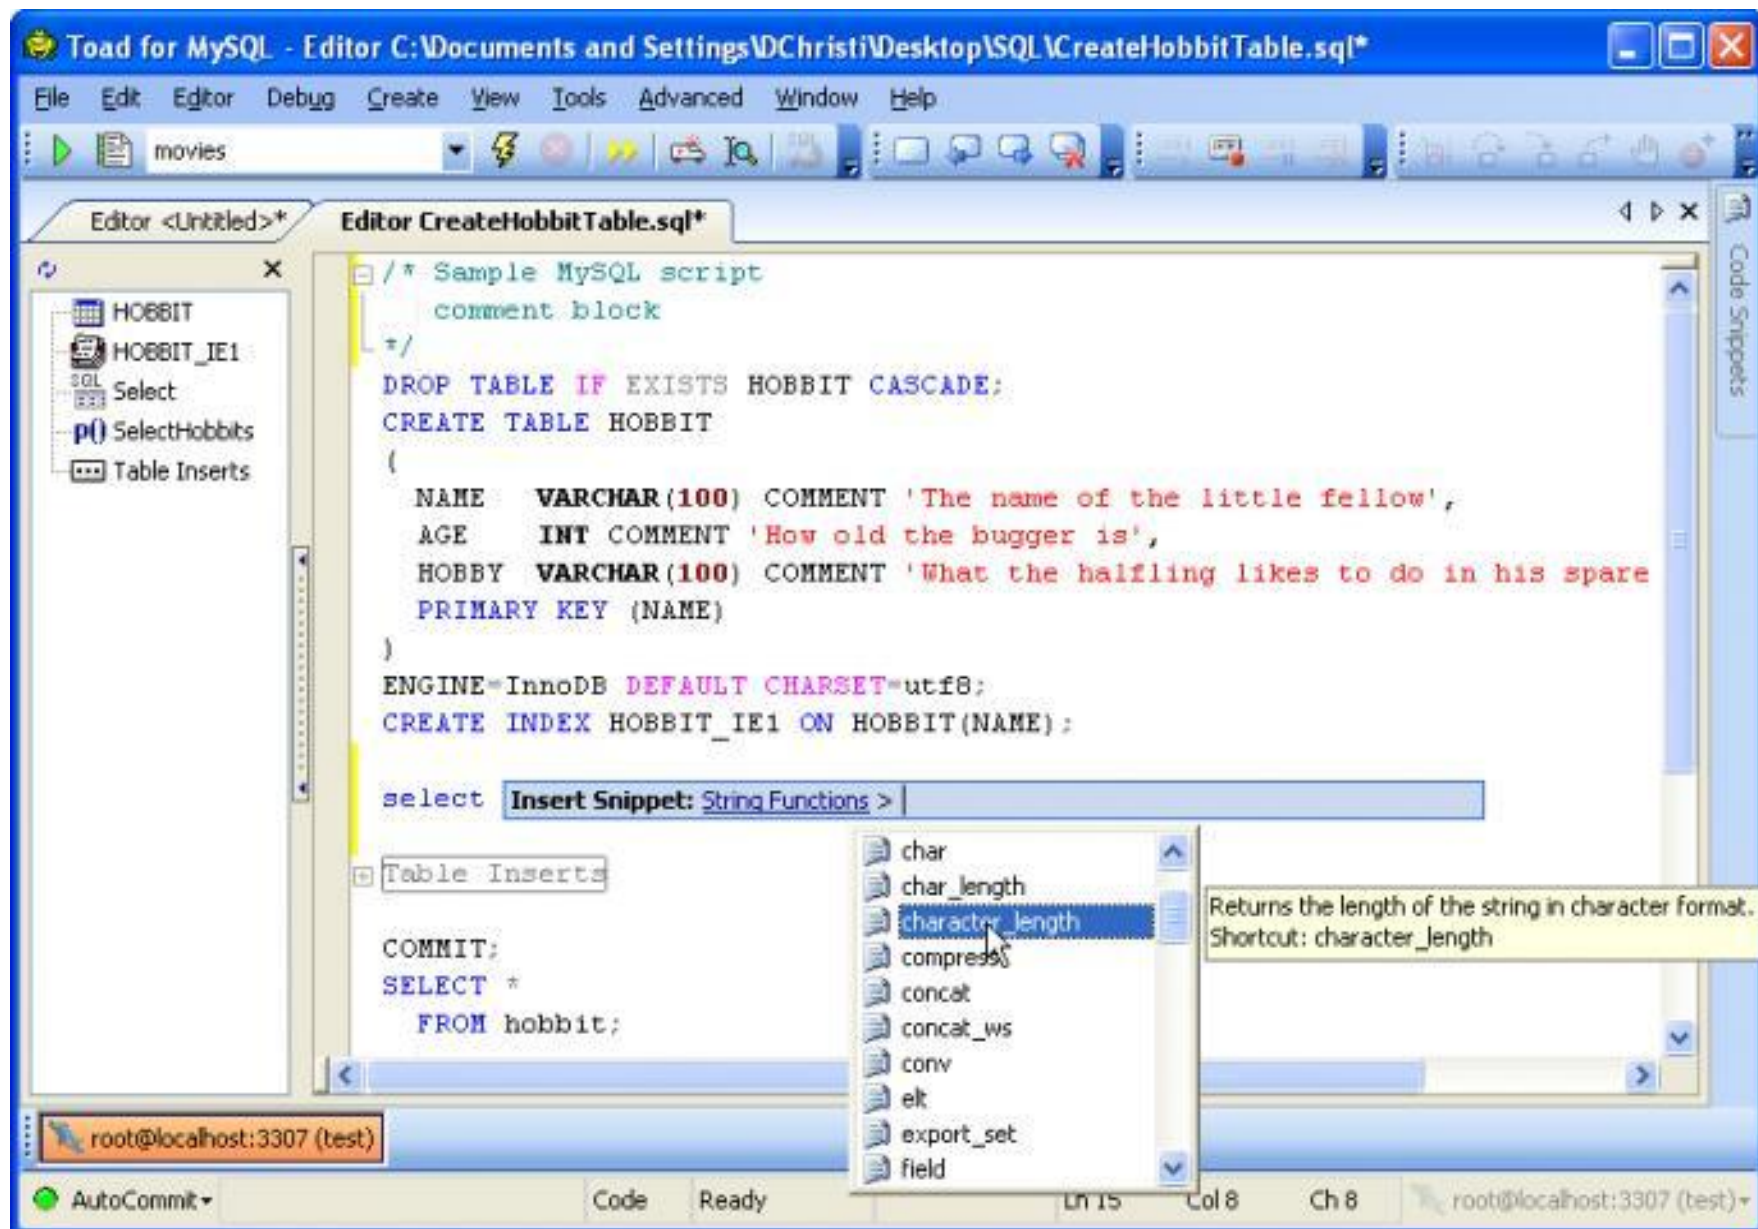
SQL>

SQL> SELECT \* FROM EXT\_EMPLOYEES WHERE BOUNUM < 10;

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE NUMBER	HIRE DATE	JOB ID	SALARY	COMMISSION PCT	MANAGER ID	DEPARTMENT ID
100	Steven	King	SKING	515.123.4567	17-JUN-87	AD_PRES	24000			
101	Neena	Kochhar	NKOCHHAR	515.123.4568	21-SEP-89	AD_UP	17000			100

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE NUMBER	HIRE DATE	JOB ID	SALARY	COMMISSION_PCT	MANAGER_ID	DEPARTMENT_ID
-------------	------------	-----------	-------	--------------	-----------	--------	--------	----------------	------------	---------------

SQL>



## **CONSTRAINTS:**

There are three types of key constraints that are most common.

Primary Key constraint

Foreign Key constraint

Unique Key constraint

Many tables will have a primary key constraint and a table may only have one primary key constraint. The primary key is one or more columns (typically one) that uniquely identifies the row. Typical examples would be an employee id for an employee table, a Stock Keeping Unit (SKU) for a product table, an order number for an orders table, and so on.

The DBMS will prevent a row from being inserted that results in any duplicate primary key— this is the ‘constraint’ part.

Foreign key constraints link two tables together. For instance, an Order table might have a customer id, that links the order to the customer id primary key in the customer table. The Order table’s primary key is the order number, but one of its foreign keys is the customer id. Foreign keys are most often used to join two tables — from a foreign key in one table to a primary key in another table. The DBMS will prevent a row from being inserted if the foreign key column has no matching row with same key in the primary key column of the parent table. For example, you can’t insert a row in the orders table with a customer number of 12345 unless there is a row in the customer table with customer id 12345.

This is the constraint that must be satisfied. Similarly, you can't delete the customer row 12345 if there are any outstanding rows in the order table with that customer id. However, you can define the foreign key constraint to 'cascade' such that deleting the customer row will 'cascade' and delete the corresponding order rows. Finally, a unique key constraint is much like a primary key constraint, but just is not the 'knighted' primary key by which rows are primarily identified. For example, a customer id may be used as the primary key for the customer table, but you might also store the social security number and having duplicate social security numbers would be a red flag that someone may have stolen an identity. By making the SSN a unique primary key constraint, the application will prevent the new customer from being added to the database.